

RAJALAKSHMI ENGINEERING COLLEGE
RAJALAKSHMI NAGAR, THANDALAM – 602 105



AI23331 Fundamentals of Machine Learning

Laboratory Record Note Book

Name :HARI G.....

Year / Branch / Section :IIIrd Year CSE - B.....

University Register No. :2116230701098.....

College Roll No. :230701098.....

Semester :VIth.....

Academic Year :2025-2026.....

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

AI23331 Fundamentals of Machine Learning

NAME	HARI G
ROLL NO	230701098
YEAR	IIIrd Year
SEC	B



RAJALAKSHMI ENGINEERING COLLEGE
An Autonomous Institution

BONAFIDE CERTIFICATE

Name:HARI G..... Academic

Year: ...2025-26... Semester: ...VIth ... Branch: ...B.E. CSE...

Register No.

230701098

*Certified that this is the bonafide record of work done by the above student in
the.....Fundamentals of Machine Learning Laboratory
during the academic year 2025 – 2026.*

Signature of Faculty in-charge

Submitted for the Practical Examination held on

Internal Examiner

External Examiner

INDEX

S.NO	Experiment	Date	GitHub
1a	A python program to implement Data Cleaning.	7/2/2026	
1b	A python program to implement univariate regression, bivariate regression and multivariate regression.	21/2/2026	
2	A python program to implement Simple linear regression using Least Square Method.	28/2/2026	
3	A python program to implement logistic model.	28/2/2026	
4	A python program to implement single layer perceptron.	7/3/2026	
5	A python program to implement multi-layer perceptron with back propagation.	7/3/2026	
6	A python program to do face recognition using SVM classifier.	14/3/2026	
7	A python program to implement decision tree.	14/3/2026	
8	A python program to implement boosting.	11/4/2026	
9a	A python program to implement KNN.	11/4/2026	
9b	A python program to implement K-means.	11/4/2026	
10	A python program to implement dimensionality reduction – PCA.	18/4/2026	

Experiment 1A: Univariate, Bivariate and Multivariate Regression

Aim	To implement univariate, bivariate and multivariate linear regression using the Iris dataset and visualize the regression results.
Dataset	Iris Dataset (from sklearn.datasets)
Tools / Libraries	Python, NumPy, Pandas, Matplotlib, scikit-learn

Program:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
from mpl_toolkits.mplot3d import Axes3D
from sklearn.datasets import load_iris

iris = load_iris()
data = pd.DataFrame(iris.data, columns=iris.feature_names)
data['target'] = iris.target

y = data['sepal length (cm)']

# ----- UNIVARIATE REGRESSION -----
X_uni = data[['petal length (cm)']]
model_uni = LinearRegression()
model_uni.fit(X_uni, y)
y_pred_uni = model_uni.predict(X_uni)
r2_uni = r2_score(y, y_pred_uni)

# Plot Univariate
plt.figure()
plt.scatter(X_uni, y, color='blue', label='Actual Data')
plt.plot(X_uni, y_pred_uni, color='red', label='Regression Line')
plt.xlabel('Petal Length (cm)')
plt.ylabel('Sepal Length (cm)')
plt.title('Univariate Regression')
plt.legend()
plt.show()

print("UNIVARIATE REGRESSION")
print("Coefficient:", model_uni.coef_)
print("Intercept:", model_uni.intercept_)
print("R-squared:", r2_uni)

# ----- BIVARIATE REGRESSION -----
X_bi = data[['petal length (cm)', 'petal width (cm)']]
model_bi = LinearRegression()
model_bi.fit(X_bi, y)
y_pred_bi = model_bi.predict(X_bi)
r2_bi = r2_score(y, y_pred_bi)

# Plot Bivariate
fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')
ax.scatter(X_bi.iloc[:, 0], X_bi.iloc[:, 1], y, color='blue')
ax.set_xlabel('Petal Length (cm)')
```

```

ax.set_ylabel('Petal Width (cm)')
ax.set_zlabel('Sepal Length (cm)')
ax.set_title('Bivariate Regression')
plt.show()

print("\nBIVARIATE REGRESSION")
print("Coefficients:", model_bi.coef_)
print("Intercept:", model_bi.intercept_)
print("R-squared:", r2_bi)

# ----- MULTIVARIATE REGRESSION -----
X_multi = data.drop(columns=['target'])
model_multi = LinearRegression()
model_multi.fit(X_multi, y)
y_pred_multi = model_multi.predict(X_multi)
r2_multi = r2_score(y, y_pred_multi)


# Plot Multivariate
plt.figure()
plt.scatter(y, y_pred_multi, color='green')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('Multivariate Regression')
plt.show()

print("\nMULTIVARIATE REGRESSION")
print("Coefficients:", model_multi.coef_)
print("Intercept:", model_multi.intercept_)
print("R-squared:", r2_multi)


```

Result / Output:


Output:



Bivariate Regression



Multivariate Regression



Experiment 1B: Data Cleaning and Preprocessing

Aim	To perform data cleaning tasks including handling missing values, detecting outliers using IQR, computing correlation, and normalizing/standardizing features.
Dataset	Custom CSV dataset (Ex1b.csv) — Heart/Cholesterol data
Tools / Libraries	Python, NumPy, Pandas, Matplotlib, Seaborn, scikit-learn

Program:

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler, StandardScaler
import seaborn as sns
import matplotlib.pyplot as plt

df = pd.read_csv("Ex1b.csv")

# Explore the dataset
df.head()
df.tail()
df.info()
df.isnull().sum()
df.describe()

# Boxplot for outlier detection
numeric_df = df.select_dtypes(include='number')
fig, axs = plt.subplots(len(numeric_df.columns), 1, figsize=(7, 18), dpi=95)
for i, col in enumerate(numeric_df.columns):
    axs[i].boxplot(numeric_df[col].dropna(), vert=False)
    axs[i].set_ylabel(col)
plt.tight_layout()
plt.show()

# Remove outliers from Cholesterol using IQR
q1, q3 = np.percentile(df["Cholesterol"].dropna(), [25, 75])
iqr = q3 - q1
lower = q1 - 1.5 * iqr
upper = q3 + 1.5 * iqr
clean_df = df[(df["Cholesterol"] >= lower) & (df["Cholesterol"] <= upper)]



# Correlation heatmap
corr = df.select_dtypes(include='number').corr()
plt.figure(dpi=130)
sns.heatmap(corr, annot=True, fmt='.2f', cmap='coolwarm')
plt.show()

# Min-Max Normalization
X = df.select_dtypes(include='number')
X_normalized = MinMaxScaler().fit_transform(X)
print(X_normalized[:5])

# Standardization
X_standardized = StandardScaler().fit_transform(X)
print(X_standardized[:5])
```


Result / Output:

Output:



Experiment 2: Simple Linear Regression using Least Square Method

Aim	To implement simple linear regression using the least square method and evaluate model performance using R2 score and Mean Squared Error.
Dataset	Custom CSV dataset (Ex2.csv) — House price data
Tools / Libraries	Python, NumPy, Pandas, Matplotlib, scikit-learn

Program:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_squared_error

df = pd.read_csv("Ex2.csv")
df.head()

x = df[['sqft_living']]
y = df['price']

x_train, x_test, y_train, y_test = train_test_split(
    x, y, test_size=0.2, random_state=42
)

model = LinearRegression()
model.fit(x_train, y_train)

y_pred = model.predict(x_test)

print("Slope (m):", model.coef_[0])
print("Intercept (c):", model.intercept_)
print("R2 Score:", r2_score(y_test, y_pred))
print("Mean Squared Error:", mean_squared_error(y_test, y_pred))

plt.scatter(x_test, y_test)
plt.plot(x_test, y_pred)
plt.xlabel("Square Feet Living")
plt.ylabel("Price")
plt.title("Simple Linear Regression: Price vs Sqft Living")
plt.show()
```

Result / Output:**Output:**

Experiment 3: Logistic Regression

Aim	To implement logistic regression for multi-class classification on the Digits dataset and evaluate the model using accuracy score and classification report.
Dataset	Digits Dataset (from sklearn.datasets)
Tools / Libraries	Python, scikit-learn

Program:

```
from sklearn.model_selection import train_test_split
from sklearn import datasets
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report

digits = datasets.load_digits()
x = digits.data
y = digits.target

x_train, x_test, y_train, y_test = train_test_split(
    x, y, test_size=0.2, random_state=42
)

model = LogisticRegression(max_iter=10000)
model.fit(x_train, y_train)

y_pred = model.predict(x_test)

accuracy = accuracy_score(y_test, y_pred)
print("Logistic Regression Accuracy:", accuracy)

print("Classification Report: ")
print(classification_report(y_test, y_pred))
```

Result / Output:

Output:

Logistic Regression Accuracy: 0.975

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	33
1	0.97	1.00	0.98	28
2	1.00	1.00	1.00	33
3	0.97	0.97	0.97	34

4	1.00	0.98	0.99	46
5	0.92	0.96	0.94	47
6	0.97	0.97	0.97	35
7	1.00	0.97	0.99	34
8	0.97	0.97	0.97	30
9	0.97	0.95	0.96	40
accuracy			0.97	360
macro avg	0.98	0.98	0.98	360
weighted avg	0.98	0.97	0.98	360

Experiment 4: Single Layer Perceptron (SLP)

Aim	To implement a Single Layer Perceptron (SLP) from scratch to learn the AND gate logic using the step activation function and weight update rule.
Dataset	AND gate truth table (manually defined)
Tools / Libraries	Python, NumPy

Program:

```
import numpy as np

# Input dataset (AND gate)
X = np.array([
    [0, 0],
    [0, 1],
    [1, 0],
    [1, 1]
])

# Target output
y = np.array([0, 0, 0, 1])

# Initialize weights and bias
weights = np.zeros(X.shape[1])
bias = 0
learning_rate = 0.2
epochs = 5

# Activation function (Step Function)
def step_function(x):
    if x >= 0:
        return 1
    else:
        return 0

# Training the perceptron
for epoch in range(epochs):
    print(f"\nEpoch {epoch+1}")
    for i in range(len(X)):
        linear_output = np.dot(X[i], weights) + bias
        y_pred = step_function(linear_output)
        error = y[i] - y_pred
        weights = weights + learning_rate * error * X[i]
        bias = bias + learning_rate * error
        print(f"Input: {X[i]}, Predicted: {y_pred}, Error: {error}")

print("\nFinal Weights:", weights)
print("Final Bias:", bias)

print("\nTesting Perceptron")
for i in range(len(X)):
    linear_output = np.dot(X[i], weights) + bias
    y_pred = step_function(linear_output)
    print(f"Input: {X[i]}, Output: {y_pred}")
```

Result / Output:

Output:

Epoch 1

Input: [0 0], Predicted: 1, Error: -1

Input: [0 1], Predicted: 0, Error: 0

Input: [1 0], Predicted: 0, Error: 0

Input: [1 1], Predicted: 0, Error: 1

Epoch 2

Input: [0 0], Predicted: 1, Error: -1

Input: [0 1], Predicted: 1, Error: -1

Input: [1 0], Predicted: 0, Error: 0

Input: [1 1], Predicted: 0, Error: 1

Epoch 3

Input: [0 0], Predicted: 0, Error: 0

Input: [0 1], Predicted: 1, Error: -1

Input: [1 0], Predicted: 1, Error: -1

Input: [1 1], Predicted: 0, Error: 1

Epoch 4

Input: [0 0], Predicted: 0, Error: 0

Input: [0 1], Predicted: 0, Error: 0

Input: [1 0], Predicted: 0, Error: 0

Input: [1 1], Predicted: 1, Error: 0

Epoch 5

Input: [0 0], Predicted: 0, Error: 0

Input: [0 1], Predicted: 0, Error: 0

Input: [1 0], Predicted: 0, Error: 0

Input: [1 1], Predicted: 1, Error: 0

Final Weights: [0.4 0.2]

Final Bias: -0.40000000000000001

Testing Perceptron

Input: [0 0], Output: 0

Input: [0 1], Output: 0

Input: [1 0], Output: 0

Input: [1 1], Output: 1

Experiment 5: Multi-Layer Perceptron (MLP) with Backpropagation

Aim	To implement a Multi-Layer Perceptron (MLP) with backpropagation from scratch to solve the XOR problem using the sigmoid activation function.
Dataset	XOR gate truth table (manually defined)
Tools / Libraries	Python, NumPy

Program:

```
import numpy as np

# Input dataset (XOR gate)
X = np.array([
    [0, 0],
    [0, 1],
    [1, 0],
    [1, 1]
])

# Target output
y = np.array([[0], [1], [1], [0]])

# Sigmoid activation function
def sigmoid(x):
    return 1 / (1 + np.exp(-x))

# Derivative of sigmoid
def sigmoid_derivative(x):
    return x * (1 - x)

# Initialize weights randomly
np.random.seed(1)
input_layer_neurons = 2
hidden_layer_neurons = 2
output_neurons = 1

W1 = np.random.uniform(size=(input_layer_neurons, hidden_layer_neurons))
W2 = np.random.uniform(size=(hidden_layer_neurons, output_neurons))

learning_rate = 0.1
epochs = 10000

for i in range(epochs):
    # Forward Propagation
    hidden_input = np.dot(X, W1)
    hidden_output = sigmoid(hidden_input)
    final_input = np.dot(hidden_output, W2)
    predicted_output = sigmoid(final_input)

    # Error
    error = y - predicted_output

    # Backpropagation
    d_predicted_output = error * sigmoid_derivative(predicted_output)
    error_hidden_layer = d_predicted_output.dot(W2.T)
    d_hidden_layer = error_hidden_layer * sigmoid_derivative(hidden_output)
```

```
# Updating weights
W2 += hidden_output.T.dot(d_predicted_output) * learning_rate
W1 += X.T.dot(d_hidden_layer) * learning_rate

print("Predicted Output:")
print(predicted_output)
```

Result / Output:

Output:

Predicted Output:

[[0.25840124]

[0.69139438]

[0.69130801]

[0.39872979]]

Experiment 6: Face Recognition using SVM Classifier

Aim	To implement a real-time face recognition system using face encodings and a Support Vector Machine (SVM) classifier trained on a custom dataset.
Dataset	Custom face image dataset (organized by person in folders)
Tools / Libraries	Python, OpenCV, face_recognition, scikit-learn (SVM)

Program:

```
import os
import cv2
import face_recognition
from sklearn import svm

# Path to dataset folder
dataset_path = r"C:\path\to\dataset"

encodings = []
names = []

# Read all person folders
for person in os.listdir(dataset_path):
    person_path = os.path.join(dataset_path, person)
    if not os.path.isdir(person_path):
        continue

    for image_name in os.listdir(person_path):
        if not image_name.lower().endswith(('.jpg', '.jpeg', '.png')):
            continue
        image_path = os.path.join(person_path, image_name)

        try:
            image = face_recognition.load_image_file(image_path)
            face_locations = face_recognition.face_locations(image)
            if len(face_locations) != 1:
                continue
            encoding = face_recognition.face_encodings(image, face_locations)[0]
            encodings.append(encoding)
            names.append(person)
        except:
            continue

# Train SVM classifier
clf = svm.SVC(gamma='scale', probability=True)
clf.fit(encodings, names)
print("Training Complete!")

# Start webcam
video_capture = cv2.VideoCapture(0)
while True:
    ret, frame = video_capture.read()
    if not ret:
        break

    rgb_frame = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
    face_locations = face_recognition.face_locations(rgb_frame)
    face_encodings = face_recognition.face_encodings(rgb_frame, face_locations)

    for (top, right, bottom, left), face_encoding in zip(face_locations, face_encodings):
```

```
name = "Unknown"
try:
    prediction = clf.predict([face_encoding])
    probabilities = clf.predict_proba([face_encoding])
    confidence = max(probabilities[0])
    if confidence > 0.7:
        name = prediction[0]
except:
    pass

cv2.rectangle(frame, (left, top), (right, bottom), (0, 255, 0), 2)
cv2.putText(frame, name, (left, top - 10),
            cv2.FONT_HERSHEY_SIMPLEX, 0.8, (0, 255, 0), 2)

cv2.imshow("Face Recognition using SVM", frame)
if cv2.waitKey(1) == 27:
    break

video_capture.release()
cv2.destroyAllWindows()
```

Result / Output:

Output:

Experiment 7: Decision Tree Classifier

Aim	To implement a Decision Tree classifier using the entropy criterion on the Iris dataset and visualize the resulting decision tree.
Dataset	Iris Dataset (from sklearn.datasets)
Tools / Libraries	Python, scikit-learn, Matplotlib

Program:

```
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt

# Load Iris dataset
iris = datasets.load_iris()
X = iris.data
y = iris.target

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42
)

# Create and train Decision Tree model
model = DecisionTreeClassifier(criterion='entropy')
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print("Predicted values:", y_pred)
print("Actual values   :", y_test)

accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Visualize the Tree
plt.figure(figsize=(12, 8))
plot_tree(model, feature_names=iris.feature_names,
          class_names=iris.target_names, filled=True)
plt.show()
```


Result / Output:

Output:

Predicted values: [1 0 2 1 1 0 1 2 1 1 2 0 0 0 0 1 2 1 1 2 0 2 0 2 2
2 2 2 0 0 0 0 1 0 0 2 1 0 0 0 2 1 1 0 0]

Actual values: [1 0 2 1 1 0 1 2 1 1 2 0 0 0 0 1 2 1 1 2 0 2 0 2 2 2
2 0 0 0 0 1 0 0 2 1 0 0 0 2 1 1 0 0]

Accuracy: 1.0



Experiment 8A: Boosting – AdaBoost

Aim	To implement the AdaBoost ensemble classifier on the Wine dataset and evaluate performance using accuracy score and classification report.
Dataset	Wine Dataset (from sklearn.datasets)
Tools / Libraries	Python, scikit-learn

Program:

```
from sklearn.model_selection import train_test_split
from sklearn.ensemble import AdaBoostClassifier
from sklearn.metrics import accuracy_score, classification_report
from sklearn.datasets import load_wine

data = load_wine()
X = data.data
y = data.target

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42
)

# Create AdaBoost model
model = AdaBoostClassifier(
    n_estimators=50,
    learning_rate=1.0,
    random_state=42
)

model.fit(X_train, y_train)
y_pred = model.predict(X_test)

print("Actual Values : ", y_test)
print("Predicted Values:", y_pred)

accuracy = accuracy_score(y_test, y_pred)
print("\nAccuracy:", accuracy)

print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

Result / Output:

Output:

Actual Values : [0 0 2 0 1 0 1 2 1 2 0 2 0 1 0 1 1 1 0 1 0 1 1 2 2 2
1 1 1 0 0 1 2 0 0 0 2 2 1 2 0 1 1 1 2 0 1 1 2 0 1 0 0 2]

Predicted Values: [0 0 2 0 1 0 1 2 1 2 0 1 0 1 0 1 1 1 0 1 0 1 1 2 2
2 1 1 1 0 0 1 2 0 0 0 2 2 1 2 0 1 1 1 2 0 1 1 2 0 1 0 0 2]

Accuracy: 0.9814814814814815

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	19
1	0.95	1.00	0.98	21
2	1.00	0.93	0.96	14
accuracy			0.98	54
macro avg	0.98	0.98	0.98	54
weighted avg	0.98	0.98	0.98	54

Experiment 9A: K-Nearest Neighbors (KNN)

Aim	To implement the K-Nearest Neighbors classifier on a custom dataset, evaluate using confusion matrix and accuracy, with feature scaling applied.
Dataset	Custom CSV dataset (knn-dataset.csv)
Tools / Libraries	Python, Pandas, scikit-learn

Program:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, accuracy_score

df = pd.read_csv("knn-dataset.csv")

X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.25, random_state=0
)

# Feature Scaling
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Create KNN model
knn = KNeighborsClassifier(n_neighbors=5, metric='minkowski', p=2)
knn.fit(X_train, y_train)

y_pred = knn.predict(X_test)

cm = confusion_matrix(y_test, y_pred)
accuracy = accuracy_score(y_test, y_pred)

print("Predicted values:")
print(y_pred)
print("\nActual values:")
print(y_test)
print("\nConfusion Matrix:")
print(cm)
print("\nAccuracy:", accuracy)
```

Result / Output:

Output:

```
[0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 1 0
0 0 0 0 0 1 0 0 0 0 1 0 0 1 0 1 1 0 0 0 1 1 0 0 1 0 0 1 0 1 0 1 0 0 0
0 1 0 0 1]
```

```
0 0 0 0 1 1 1 0 0 0 1 1 0 1 1 0 0 1 0 0 0 1 0 1 1 1]
```

Confusion Matrix:

```
[[64  4]
```

```
 [ 3 29]]
```

Accuracy: 0.93

Experiment 9B: K-Means Clustering

Aim	To implement K-Means clustering and determine the optimal number of clusters using the Silhouette Score method.
Dataset	Custom CSV dataset (kmeans-dataset.csv) — Mall Customers data
Tools / Libraries	Python, Pandas, Matplotlib, scikit-learn

Program:

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score

df = pd.read_csv("kmeans-dataset.csv")

# Select Annual Income and Spending Score columns
X = df.iloc[:, [3, 4]].values

# Compute silhouette scores for different values of k
silhouette_scores = []

for k in range(2, 11):
    kmeans = KMeans(n_clusters=k, init='k-means++', random_state=0)
    labels = kmeans.fit_predict(X)
    score = silhouette_score(X, labels)
    silhouette_scores.append(score)
    print(f"k = {k}, Silhouette Score = {score:.4f}")

# Plot the silhouette scores
plt.plot(range(2, 11), silhouette_scores, marker='o')
plt.title("Silhouette Method for Optimal k")
plt.xlabel("Number of Clusters")
plt.ylabel("Silhouette Score")
plt.show()

best_k = range(2, 11)[silhouette_scores.index(max(silhouette_scores))]
print("\nBest number of clusters:", best_k)
```

Result / Output:


Output:

```
k = 2, Silhouette Score = 0.3849
k = 3, Silhouette Score = 0.4676
k = 4, Silhouette Score = 0.4932
k = 5, Silhouette Score = 0.5539
k = 6, Silhouette Score = 0.4571
k = 7, Silhouette Score = 0.4549
```

$k = 8$, Silhouette Score = 0.4607

$k = 9$, Silhouette Score = 0.4579

$k = 10$, Silhouette Score = 0.4222



Experiment 10: Dimensionality Reduction – PCA

Aim	To apply Principal Component Analysis (PCA) on the Iris dataset to reduce dimensionality from 4 features to 2 principal components and analyze the explained variance.
Dataset	Iris Dataset (from sklearn.datasets)
Tools / Libraries	Python, Pandas, scikit-learn

Program:

```
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
import pandas as pd

iris = load_iris()
X = iris.data
y = iris.target

# Standardize the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Apply PCA to reduce from 4 features to 2 principal components
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

# Create a DataFrame to view the result
df = pd.DataFrame(
    data=X_pca,
    columns=['Principal Component 1', 'Principal Component 2']
)
df['Species'] = y

print(df.head())

print("\nExplained Variance Ratio:")
print(pca.explained_variance_ratio_)
```

Result / Output:

Output:

	Principal Component 1	Principal Component 2	Species
0	-2.264703	0.480027	0
1	-2.080961	-0.674134	0
2	-2.364229	-0.341908	0
3	-2.299384	-0.597395	0
4	-2.389842	0.646835	0

